

Report ASP20 Boost

First Report

Johannes Strauß

Levin Wiebelt

Sebastian Aristizabal

04 Juni 2020

Contents

1. Introduction:	2
2. Concept: Boosting	3
Location-Booster	3
Scale-Booster	3
3. Implementation	4
4. Functionality: Componentwise Boosting	5
5. Prospects	6
Further Functionalities	6
Stability	6
Performance	6
Further	6
Application	6

1. Introduction:

The course “Advanced Statistical Programming with R” consists in implementing a package to address location-scale regression. The sample dataset consists of a response y , that’s expectation is in linear relationship to a set of predictors x_i . Its variance, in contrast, is dependent on a different set of linear predictors z_i , however transformed by the response function $g(x) = \exp(x)$. This ensures positiveness of the variances. The goal approach is to estimate the effects of both: x_i on the expectation of y (location) - this is captured by the estimates β , and z_i on the variance of y (scale) - this is captured by the estimates γ . Our student group **asp20boost** solves this task applying the concept of boosting.

In part 2 of this report we explain the concept of boosting and elaborate on how it suites to address the problem of location-scale regression. In the following part 3 we describe our implementation milestones, whereby part 4 especially elaborates on componentwise boosting. In part 5 we present our thoughts on upcoming challenges and critical aspects of the package development.

2. Concept: Boosting

The first step in our progress was to understand the theoretical concept of boosting, and which types of problems it is able to solve. In the following the term “booster” refers to a boosting algorithm.

Location-Booster

To explain the concept and follow our process of understanding we first explain boosting for location parameters. Boosting scale parameters requires another way of thinking about it and constituted a first milestone in our understanding and implementation. The boosting concept relies on the idea of iterative estimation. To estimate the location parameters β via boosting, one starts with initial estimates, which may be far from optimal. However, with estimates and the response vector at hand, residuals can be calculated. The boosting-algorithm focuses on those calculated residuals. The effect of the predictors x_i on these residuals is estimated by least squares and the resulting effect is added to the location-estimate β , adjusted by some learning rate. This yields a better fit of the location-model, which results in smaller residuals. These new residuals are estimated again in the next iteration, yielding smaller effect sizes, and hence a convergence of the boosting algorithm towards the OLS-estimate. A simple booster stops after a fixed number of iterations is reached. (**comp. Kneib et al**)

Scale-Booster

[Loss-function erklären: Boosting can be viewed as lowering iteratively the amount of a loss function, which is in our case the score of the normal distribution]

The concept for boosting scale parameters is the same. However, residuals for the variance estimates can not be calculated, since there is no such thing as “observed” variances for individual response values. The solution is to substitute the calculated residuals with a different term: the derivative of the loss function, which in our case is the log-likelihood of the normal distribution. Instead of estimating stepwise residuals, one estimates hence stepwise the score.

The resulting loss function then, is defined by sum of squares of the term: u_i (=deviance-residuals) - inverse predicted scores * working-residuals

We employed the resid-function of the `asp20model` class, on which our class is built on.

3. Implementation

In the current version of the `asp20boost`-package, the boosting algorithm is already implemented in order to allow componentwise boosting, which will be explained in part 4. For good explanation and for following our development process in this part we first describe the implementation of ordinary boosting - which however does not appear in our code anymore.

A simple booster for the location parameters β works without extending the `LocationScaleRegression-Class`. The current residuals may be extracted with the `resid()`-method, then estimated. The location parameter is updated. This causes the `LocationScaleRegression-Class` to calculate updated state-dependent objects, such as the loglikelihood, gradients and the residuals. This is what's needed to repeat the procedure of estimating residuals and updating location parameters.

The simple booster for γ consists in the same two repeating steps of first estimating a term, and second updating the scale parameters γ by adding the estimated effect - adjusted by a learning rate. The term to be estimated, however, is not the residual, but the derivative of the loglikelihood of the normal distribution. To calculate this, our code makes use of the `resid-function()`, but this time passing the argument "deviance", which results in residuals adjusted by the fitted scale estimates.

[Is the gamma booster dependent on simultaneous running of beta booster?]



Having implemented a boosting algorithm for location and scale, the code-basis for our package was achieved.

4. Functionality: Componentwise Boosting

A usefull functionality of boosting is componentwise boosting. The idea is to not update a whole parameter vector, but only one entry of it. The entry chosen for the update at this juncture is the one yielding the best improvement in the sense of lowering the loss function. The loss function in our case is the loglikelihood of the normally distributed data. Hence in each iteration only one component of the location-, as well as one component of the scale-parameters are estimated.

The implementation in our package works via the extension of the `LocationScaleRegressionClass` by two active fields “`bestFittingVariableBeta`” and “`bestFittingVariableGamma`”. These functions partition the design matrix X , respectively Z , into its single columns and then estimate the residuals - respectively the score - seperately for each component, and determine loss functions for a hypothetical update with the respective component. Comparing with the old loss function value, the highest loss-improvement may be determined and the respective component is used to update the parameter-vector.

We are in the process of reconsidering the design of this implementation. Other design possibilities are the following:

- create public fields for the heavily used score-function values
- move the calculation of componentwise losses to an external function
- harmonize boosting and componentwise boosting into one external function, determining the mode of operation by an argument ‘`componentwise = TRUE`’

Another conceptual questions that comes up is if the best loss-improvement may be indicated by the already existing gradients, and hence there is no need for extra calculation.

5. Prospects

Further Functionalities

The most important functionality we intend to implement in our package is variable selection via cross validation. Working out the theoretical concept behind this idea is one of our next milestones.

One further, rather loose, idea is to find ways to optimize the learning rate in our boosting algorithm.

Stability

We intend to implement further automated unit tests. This will enable us to assess quickly the stability of our code given differing inputs. Departing from this point we intend to reduce proneness to input errors.

Performance

- A major performance problem in our code is that small learning rates for the beta-boosters cause processing time of the code to increase sharply. This remains to be solved.

*A good choice of starting values for beta and gamma may enhance performance. A possible candidate is the mean of the response.

Further

- Allow user-input of the LocationScaleRegressionBoost-model in form of a dataframe, for example by allowing to pass an optional “data”-argument as known from the lm-call. This may reduce proneness to input-errors.
- Visualize the results of the boosted estimates
- Document the extension of the R6-Class properly and inherit documentation of the LocationScaleRegression-Class.



Application

[Johannes: Munich Rent Data? - Identify Outliers]